



```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

```
In [2]: df = pd.read_csv("creditcard.csv")
```

```
In [3]: df.head()
```

```
Out[3]:
```

	Time	V1	V2	V3	V4	V5	V6	V7
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941

5 rows × 31 columns

```
In [4]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Time        284807 non-null float64
1   V1          284807 non-null float64
2   V2          284807 non-null float64
3   V3          284807 non-null float64
4   V4          284807 non-null float64
5   V5          284807 non-null float64
6   V6          284807 non-null float64
7   V7          284807 non-null float64
8   V8          284807 non-null float64
9   V9          284807 non-null float64
10  V10         284807 non-null float64
11  V11         284807 non-null float64
12  V12         284807 non-null float64
13  V13         284807 non-null float64
14  V14         284807 non-null float64
15  V15         284807 non-null float64
16  V16         284807 non-null float64
17  V17         284807 non-null float64
18  V18         284807 non-null float64
19  V19         284807 non-null float64
20  V20         284807 non-null float64
21  V21         284807 non-null float64
22  V22         284807 non-null float64
23  V23         284807 non-null float64
24  V24         284807 non-null float64
25  V25         284807 non-null float64
26  V26         284807 non-null float64
27  V27         284807 non-null float64
28  V28         284807 non-null float64
29  Amount      284807 non-null float64
30  Class       284807 non-null int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB

```

```
In [5]: df.describe()
```

Out[5]:

	Time	V1	V2	V3	V
<b>count</b>	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
<b>mean</b>	94813.859575	1.175161e-15	3.384974e-16	-1.379537e-15	2.094852e-15
<b>std</b>	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00
<b>min</b>	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+01
<b>25%</b>	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01
<b>50%</b>	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-01
<b>75%</b>	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01
<b>max</b>	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01


8 rows × 31 columns

```
In [6]: df["Class"].value_counts()
```

```
Out[6]: Class
0      284315
1         492
Name: count, dtype: int64
```

```
In [7]: sns.heatmap(df.corr(), cmap='coolwarm')
```

```
Out[7]: <Axes: >
```



```
In [8]: df['Class'].value_counts()
```

```
Out[8]: Class
0      284315
1         492
Name: count, dtype: int64
```

```
In [9]: normal=df[df.Class==0]
fraud=df[df.Class==1]
print(normal.shape)
print(fraud.shape)
```

```
(284315, 31)
(492, 31)
```

```
In [10]: df.groupby('Class').mean()
```

```
Out[10]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10	V11	V12	V13	V14	V15	V16	V17	V18	V19	V20	V21	V22	V23	V24	V25	V26	V27	V28	Class
0	94838.202258	0.008258	-0.006271	0.012171	-0.007860	0.005453	0.0024	0.0015	0.0015	0.0015	0.0015	0.0015	0.0015	0.0015	0.0015	0.0015	0.0015	0.0015	0.0015	0.0015	0.0015	0.0015	0.0015	0.0015	0.0015	0.0015	0.0015	0.0015	0.0015	
1	80746.806911	-4.771948	3.623778	-7.033281	4.542029	-3.151225	-1.3977	-1.3977	-1.3977	-1.3977	-1.3977	-1.3977	-1.3977	-1.3977	-1.3977	-1.3977	-1.3977	-1.3977	-1.3977	-1.3977	-1.3977	-1.3977	-1.3977	-1.3977	-1.3977	-1.3977	-1.3977	-1.3977		

2 rows x 30 columns

```
In [11]: normal_sample=normal.sample(n=492)
new_file=pd.concat([normal_sample,fraud],axis=0)
```

```
In [12]: new_file.head(10)
```

```
Out[12]:
```

	Time	V1	V2	V3	V4	V5	V6
<b>30592</b>	35995.0	1.053277	-0.013309	0.282071	1.262380	-0.328526	-0.410608
<b>2806</b>	2364.0	-1.303390	1.451622	0.210419	-0.365147	-0.314256	-0.081985
<b>20625</b>	31164.0	1.534286	-0.938978	-0.739794	-1.907128	-0.060218	0.496395
<b>111760</b>	72358.0	-0.866577	1.577138	1.282551	2.052141	0.711365	1.136185
<b>279506</b>	168914.0	0.036777	0.825349	0.284599	-0.588419	0.401272	-1.056775
<b>13413</b>	23716.0	-1.448540	0.405678	1.549291	1.105241	0.340938	1.706729
<b>252620</b>	155892.0	2.057085	0.156833	-1.794370	1.270376	0.614585	-0.727465
<b>118117</b>	74957.0	-0.884674	0.789085	-0.262496	-2.411150	1.991090	3.151908
<b>237771</b>	149383.0	1.994414	0.198878	-1.527961	1.290176	0.440207	-0.878733
<b>232286</b>	147113.0	1.629657	-1.005061	-0.830974	0.304987	-0.524145	-0.240532

10 rows × 31 columns

```
In [13]: new_file['Class'].value_counts()
```

```
Out[13]: Class
0      492
1      492
Name: count, dtype: int64
```

```
In [14]: new_file.groupby('Class').mean()
```

```
Out[14]:
```

	Time	V1	V2	V3	V4	V5	V6
<b>Class</b>							
<b>0</b>	96173.363821	-0.115545	-0.094846	-0.034031	-0.072742	0.082600	-0.0793
<b>1</b>	80746.806911	-4.771948	3.623778	-7.033281	4.542029	-3.151225	-1.3977

2 rows × 30 columns

```
In [15]: X=new_file.drop(columns='Class',axis=1)
Y=new_file['Class']
X_train,X_test,Y_train,Y_test=train_test_split(X,Y,test_size=0.2,stratify=Y,ra
model=LogisticRegression()
model.fit(X_train,Y_train)
```

```
C:\Users\Admin\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\linear_model\_logistic.py:473: ConvergenceWarning: lbfgs failed to converge after 100 iteration(s) (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT
```

Increase the number of iterations to improve the convergence (max\_iter=100).

You might also want to scale the data as shown in:
















<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

Out[15]:

LogisticRegression		
Parameters		
	penalty	'l2'
	dual	False
	tol	0.0001
	C	1.0
	fit_intercept	True
	intercept_scaling	1
	class_weight	None
	random_state	None
	solver	'lbfgs'
	max_iter	100
	multi_class	'deprecated'
	verbose	0
	warm_start	False
	n_jobs	None
	l1_ratio	None

```
In [17]: X_train_prediction=model.predict(X_train)
training_data_acuracy=accuracy_score(X_train_prediction,Y_train)*100
print(f"Accuracy of Model: {training_data_acuracy}%")
```

Accuracy of Model: 94.79034307496823%

In [ ]: